

UNIVERSITY OF CAMBRIDGE

C1 Research Computing - Coursework Report

Jacob Tutt (JLT67)

Department of Physics, University of Cambridge

December 15, 2024

Word Count: 2998

1 Introduction

This report presents **dual_autodiff**, an automatic differentiation package, through discussing the motivation for its development, its core functionalities, and the rationale behind its structure and design. Its primary focus is to provide a critical review of the package's architecture, implementation and performance while discussing potential future developments.

1.1 Background

Automatic differentiation (AD) has been a well-established area of research in fields such as computational fluid dynamics and atmospheric sciences for many years [2]. Its popularity stems from its ability to compute derivatives of functions with machine precision while significantly reducing the computational cost compared to traditional numerical methods [1]. In recent years, its role has become increasingly important in machine learning, where derivatives, particularly gradients and Hessians, are ubiquitous in optimisation and training [2].

The mathematical framework of forward-mode AD is provided by dual numbers, which are defined as:

$$\mathbb{D} = \{a + b\epsilon : a, b \in \mathbb{R}\} \quad (1)$$

where ϵ is a non-zero number with the defining property:

$$\epsilon^2 = 0.$$

When functions are evaluated using dual numbers $a + b\epsilon$, their Taylor expansion can be expressed as:

$$f(a + b\epsilon) = f(a) + f'(a)(b\epsilon) + \frac{f''(a)}{2!}(b\epsilon)^2 + \frac{f'''(a)}{3!}(b\epsilon)^3 + \dots \quad (2)$$

Since any higher-order terms than ϵ^2 vanish by definition, the result simplifies to:

$$f(a + b\epsilon) = f(a) + bf'(a)\epsilon \quad (3)$$

Hence, the result of evaluating $f(a + b\epsilon)$ is a dual number where:

- Real Part ($f(a)$): The function evaluated at a .
- Dual Part ($bf'(a)\epsilon$): The functions first derivative evaluated at a , scaled by b .

This property forms the foundation of automatic differentiation. By evaluating a function using a dual number with $b = 1$, both the function value and its derivative are computed simultaneously. Further proofs and resources for a greater understanding are given in the packages [documentation](#). The **dual.autodiff** package aims to provide a framework to define dual numbers, evaluate them for breadth of functions within python and finally exploit this structure to provide an automatic differentiation tool.

2 Package Functionalities

2.1 Base Dual Class

At the core of this package is the **Dual** class, which defines the properties of the dual numbers, allowing their initialisation and storing its constituent parts. It also defines the formatting for representation (`__repr__`) of the dual number as a string. A demonstration of this is shown below:

Listing 1: Example initialisation of the Dual class

```

1  x = Dual(2, 1)
2  print(x.real)
3  print(x.dual)
4  print(x)

```

Output:

```

2
1
Dual(real=2.0, dual=1.0)

```

2.2 Dual Numbers' Operations

The package goes on to support a range of basic mathematical operations for dual numbers, such as the addition, multiplication and power operations. This allows the class to integrate with the python arithmetic operators (`+`, `-`, `*`, `/`, `**`) making it intuitive to use and more robust for the user when defining functions. This is achieved by overloading the python operators in the Dual class eg (`__add__`, `__mul__`, `__pow__`). The derivations of these operations in the context of the dual numbers is outlined in Table 1, for the numbers $a + b\epsilon$ and $c + d\epsilon$:

Operation	Result
Addition (+)	$(a + c) + (b + d)\epsilon$
Subtraction (-)	$(a - c) + (b - d)\epsilon$
Multiplication (*)	$ac + (ad + bc)\epsilon$
Division (/)	$\frac{a}{c} + \frac{bc - ad}{c^2}\epsilon$
Power (**)	$a^n + na^{n-1}b\epsilon$

Table 1: Dual Numbers Operations

The package additionally supports these operators when for a combination of dual numbers and scalars, and defines the reverse operations for the commutative operators (eg. `__radd__`, `__rmul__`). Without this functionality, the automatic differentiation tool (Section 2.5) would be limited to functions which do not contain scalars.

The inclusion of the equality and comparison operators was less trivial, as dual numbers are not directly comparable. This package builds on the theory provided by Cheng [3] (1994) which defines equality for dual numbers as the equivalence of both their real and dual components but goes on to make comparisons based solely on their real parts (shown in Table 2).

Operation	Result
Equality (==)	$(a == c) \wedge (b == d)$
Inequality (!=)	$(a != c) \wedge (b != d)$
Less Than (<)	$a < c$
Greater Than (>)	$a > c$
Less Than or Equal To ($\leq$)	$a \leq c$
Greater Than or Equal To ($\geq$)	$a \geq c$

Table 2: Dual Numbers Comparison

A small range of examples is provided below, with a much more extensive list available in the packages [documentation](#) and notebook.

Listing 2: Example of Dual Number Operations

```

1  x = Dual(2, 1)
2  y = Dual(3, 2)
3  print(x + y)
4  print(x * y)
5  print(x ** 2)
6  print(x == y)
7  print(x < y)

```

Output:

```

Dual(real=5.0, dual=3.0)
Dual(real=6.0, dual=7.0)
Dual(real=4.0, dual=4.0)
False
True

```

2.3 Dual Numbers' Functions

The package also defines more complex mathematical functions within the dual numbers class, a full list of which is defined in Table 3. These functions $f(a + b\epsilon)$ are implemented using the chain rule property defined in Equation 3, allowing the real and dual parts to be computed independently.

Function	Real Part	Dual Part, ϵ
$\sin(a + b\epsilon)$	$\sin(a)$	$b \cos(a)$
$\cos(a + b\epsilon)$	$\cos(a)$	$-b \sin(a)$
$\tan(a + b\epsilon)$	$\tan(a)$	$b \sec^2(a)$
$\sinh(a + b\epsilon)$	$\sinh(a)$	$b \cosh(a)$
$\cosh(a + b\epsilon)$	$\cosh(a)$	$b \sinh(a)$
$\tanh(a + b\epsilon)$	$\tanh(a)$	$b \operatorname{sech}^2(a)$
$\arcsin(a + b\epsilon)$	$\arcsin(a)$	$\frac{b}{\sqrt{1-a^2}}$
$\arccos(a + b\epsilon)$	$\arccos(a)$	$-\frac{b}{\sqrt{1-a^2}}$
$\arctan(a + b\epsilon)$	$\arctan(a)$	$\frac{b}{1+a^2}$
$\exp(a + b\epsilon)$	$\exp(a)$	$b \exp(a)$
$\log(a + b\epsilon)$	$\log(a)$	$\frac{b}{a}$
$\sqrt{a + b\epsilon}$	$\sqrt{a}$	$\frac{b}{2\sqrt{a}}$
$\operatorname{pow}(a + b\epsilon, n)$	a^n	$na^{n-1}b$

Table 3: Dual Numbers Functions

These are initially defined as member functions of the `Dual` class, and only callable using `Dual.func()`. This would result in non-intuitive definitions for compound functions being required by user. Hence, the package was extended to include these functions as independent classes within the `autodiff.tools` module. These functions are defined to be able to take scalars, dual numbers, or arrays of either (discussed further in Section 2.4). Examples of single input functions are shown below:

Listing 3: Example of Dual Number Functions

```

1  x = Dual(2, 1)
2  print(x.sin())
3  print(sin(x))
4  print(log(x))
5  print(tan(x) + cos(x) * pow(x, 2))

```

Output:

```

Dual(real=0.9093..., dual=-0.4161...)
Dual(real=0.9093..., dual=-0.4161...)
Dual(real=0.6931..., dual=0.5)
Dual(real=-3.8496..., dual=0.4726...)

```

2.4 Numpy Integration

In fields such as machine learning, it is often necessary to differentiate functions across batches or large arrays of data points, highlighted by ongoing research into the parallelisation of automatic differentiation [5]. In order to address this need (at a elementary level) and enhance the package’s utility in modern applications, Version 1.1 of `dual_autodiff` introduced seamless integration with NumPy arrays. This implementation is achieved by exploiting numpy’s ability to identify overwritten operators and functions. The implications of this approach on computational efficiency are further discussed in Section 2.7. Its practical application is shown below:

Listing 4: Example of Numpy Integration

```

1  x = np.array([Dual(2, 1), Dual(3, 2)], ...)
2  print(log(x))
3  print(tan(x) + cos(x) * pow(x, 2))

```

Output:

```

[ Dual(real=0.6931, dual=0.5000), Dual(real=1.0986, dual=0.6667), ...]
[ Dual(real=-3.8496, dual=0.4726), Dual(real=-9.0525, dual=-12.3794), ...]

```

2.5 Automatic Differentiation

Using the Dual number’s framework described in the previous sections, an automatic differentiation function (`auto_diff`) is introduced to the `autodiff_tools` module.

The automatic differentiation function takes a function (composed of the arithmetic operations and functions defined in Tables 1 and 3) and a scalar value(s) at which to evaluate it and its derivative. Due to the nature of dual numbers this can be simply implemented by:

1. Accepting a function and a scalar value (x) as input.
2. Initialising a dual number ($\mathbb{D} = x + \epsilon$) from the scalar value.
3. Evaluating the function at a $\mathbb{D}$, $f(\mathbb{D})$
4. Return Function Value: $f(a) = \text{real}(f(\mathbb{D}))$
5. Return Derivative Value: $f'(a) = \text{dual}(f(\mathbb{D}))$

An example of its application is shown below:

Listing 5: Example of `autodiff` Function

```

1  def f(x):
2      return x ** 2 + 2 * x + 1
3
4  x = 2
5  value, derivative = auto_diff(f, x)
6  print(value)
7  print(derivative)
8
9  # Similarly for an array of values
10 x_array = np.array([2, 3, 4])
11 value_array, derivative_array = auto_diff(f, x_array)
12 print(value_array)
13 print(derivative_array)

```

Output for single input:
Value: 9.0
Derivative: 6.0
Output for array of inputs:
Value: [9.0, 16.0, 25.0]
Derivative: [6.0, 8.0, 10.0]

This functionality was further expanded in the function `multi_auto_diff` to support the extended case in which the user wishes to evaluate multiple functions and their derivatives simultaneously. This can be achieved by passing a list of functions as well as the scalar values at which to evaluate them. The performance of the automatic differentiation is demonstrated below in Section 2.6.

2.6 Validation of Automatic Differentiation

The validity of the package's automatic differentiation was tested by comparing its results to the analytical derivative for a range of functions. These functions (of varying complexity) and their analytical derivatives are defined in Table 4.

Name	Function	Analytical Derivative
f_1	$x^2 + x$	$2x + 1$
f_2	$\sin(x)$	$\cos(x)$
f_3	$\log(\sin(x)) + x^2 \cos(x)$	$\frac{\cos(x)}{\sin(x)} + 2x \cos(x) - x^2 \sin(x)$

Table 4: Functions and their analytical derivatives used to verify the packages automatic differentiation ability

The functions and their analytical derivatives were evaluated over the range $x \in (0, \pi)$ and compared to the results from automatic differentiation, using the `multi_auto_diff` function, more examples of the exact implementation can be found in the packages [documentation](#) and notebook. The results of this comparison are shown in Figure 1.




Figure 1: Comparison of the analytical and automatic differentiation results for the functions defined in Table 4, showing both the function and their derivatives' values as well as the maximum difference between them.

For all functions, the automatic differentiation results were consistent with the analytical derivatives, with $f_1(x)$ and $f_2(x)$ having zero deviation from the analytical form and the maximum deviation for $f_3(x)$ being 7.11×10^{-15} , which falls within the expected numerical error for floating-point calculations. Overall, this demonstrates the package's ability to produce exact derivatives through automatic differentiation.

2.7 Comparison with Numerical Differentiation

To justify the use of automatic differentiation, the package’s performance is compared with traditional numerical methods, using accuracy, speed, and memory usage as metrics. Both forward difference (Eq 4) and central difference (Eq 5) methods are comparison against. The benchmark that these are each evaluated for is differentiating $f_3(x)$ from Table 4 for 100 data points over the range of $x \in (0, \pi)$. Each metric provided is averaged for 100 runs.

$$f'(x) \approx \frac{f(x+h) - f(x)}{h} \quad (4)$$

$$f'(x) \approx \frac{f(x+h) - f(x-h)}{2h} \quad (5)$$

Firstly, we compare the method’s computational overheads (Table 5), and can see a clear advantage in terms of speed and memory usage for the numerical approaches, with typically both numerical methods being two order of magnitude faster and using on average $\approx 78\%$ of the memory. Although initially unexpected, it can be attributed to each of the methods being performing simultaneously on 100 data points for each iteration. Hence the numerical derivatives, which fully exploit numpy’s C backend, benefit from significant performance improvements through vectorisation. Improving the package’s efficiency by further increasing in C backend usage is discussed in Section 3.

Method	Avg Run Time (s)	Avg Memory (kb)
AutoDiff	0.006733	6.10
Forward Difference ($h=0.1$)	0.000019 (0.28%)	4.34 (71%)
Forward Difference ($h = 1 \times 10^{-10}$)	0.000017 (0.25%)	4.34 (71%)
Central Difference ($h=0.1$)	0.000023 (0.34%)	5.21 (85%)
Central Difference ($h = 1 \times 10^{-10}$)	0.000019 (0.28%)	5.20 (85%)

Table 5: Comparison of Average Run Time and Memory Usage for Different Methods

The second factor to consider is the accuracy of numerical methods in comparison with automatic differentiation’s, which was shown to be exact (within machine error) in Section 2.6. Figure 2a shows the performance of the forward difference method with $h = 0.1$ as a demonstration of where the numerical methods fails to provide accurate results. For the case of $f_3(x)$, this can be seen to be as x approaches 0 and π , corresponding to where the function is discontinuous and hence has very large second derivatives. Therefore automatic differentiation shows advantages towards ill-defined regions, as it is able to provide exact derivatives for these points, whereas numerical methods are limited by the step size’s relation to the function’s second derivative.

Figure 2 shows the maximum deviation from the analytical derivative, for the forward and central difference methods, across varying step sizes. It can be seen that although the central difference method is more accurate than the forward difference method, both are significantly effected by the step size. Notably, however it suggests that there is an optimal step size, h , at which the methods reaches a maximum accuracy before the error increases again. This can be attributed to numerical integrations requirement to balance the truncation errors with numerator’s round off error (catastrophic cancellation [4]). The optimal step size for the central difference method was found to be $h = 7.85 \times 10^{-7}$ with a maximum error of 1.27×10^{-6} and similarly for the forward difference method $h = 1.44 \times 10^{-9}$ with a maximum deviation of 7.18×10^{-9} . Although important to consider that this is an extreme case, the automatic differentiation shows clear advantages in accuracy, with numerical methods limited to an error at least 6 order of magnitude greater than machine precision. Thus in application where very high degrees of accuracy are required, such as scientific computing and machine learning, automatic differentiation would be preferred.




Figure 2: Comparison of accuracy for numerical differentiation across step sizes

3 Exploting C's Optimised Performance

As highlighted in Section 2.7, the performance is constrained by Python's interpreted execution, which is inherently slower than C's compiled approach. This section discusses the steps taken to maximise the package's performance by exploiting C. Firstly, through Cythonisation and then through future developemnt of the packages use of numpy's C backend.

3.1 Cythonisation

A cythonised sub-package (`dual_autodiff_x`) was then created, a proccess in which the package is converted to C allowing it to be precompiled, thus improving its run-time performance. The structure of the package is further discussed in Section 4. The performance of the cythonised package (`dual_autodiff_x`) was then compared with the original python package (`dual_autodiff`), by evaluating the automatic derivative for a range of functions (detailed in figure 3). This was preformed for single inputs ($x = 1$) and an array of 8 inputs, each ($x = 1$) for comparability.




Figure 3: Comparison of cython and python package's computational overheads for automatic differentiation for both single and array inputs

The cythonised package shows consistent runtime decreases across all functions and input types. For single inputs, we see runtime decreases in the range of 14-29%. Although significant, the majority of the original python package's functions are implemented through the C dependent 'math' library and thus already exploiting C's advantages and resulting in a marginal increase. Notably, however the array inputs show a smaller fractional change, typically 2/3 of their single input counter parts. This helps show support for the packages success in vectorisation (section 2.4) as the packages is already exploiting numpy's C backend more, making its performance less changeable for cythonisation.

On the other hand the memory usage of the cythonised package typically shows a significant increase, upwards of 50% for complex functions such as $f_3(x)$. This is expected due to C's use of non-dynamic memory allocation, which although causes greater memory demand, is one of the features which assists the improved computational performance.




Figure 4: Scaling of automatic differentiation of the function $\sin(x)$ with array size

Finally, the scaling of run time with array size for the automatic differentiation of $\sin(x)$ is shown in Figure 4. It shows a linear trend for both the python and cythonised sub-packages, with the former showing a consistent ($\approx 10\%$) increase. Both therefore show a computational complexity of $O(n)$. Contrary to before this shows that the packages are performing the calculations within the array sequentially rather than in parallel, and thus the packages vectorisation still requires further development.

3.2 Future Developments for C

Version 1.2 will aim to further exploit C's capabilities by redefining the `__array_ufunc__` method to enable full vectorization of operations. The correct overwritten operators can be found within the `Future_Development` folder of the package's repository (note this would typically not be included in the package distribution but has been for assesment purposed). Although this method has been shown to work for well defined cases, the Dual class's initialisation requires rebuilding to support all cases robustly, so is left for future development rather than compromising the packages current functionality. A simplified example of the implementation of this method is shown below for `np.sin`:

Listing 6: Example of `__array_ufunc__` Method

```

1  def __array_ufunc__(self, ufunc, method, *inputs, **kwargs):
2      if ufunc == np.sin:
3          real_result = np.sin(real_parts)
4          dual_result = np.cos(real_parts) * dual_parts
5          return Dual(real_result, dual_result)

```

4 Development of Package

This section provides a review of the software development practises followed during the coursework and provides the motivation behind some of the decisions made.

4.1 Error Handling

Extensive error handling was implemented throughout to ensure robustness as well as clear and informative error messages making the package more user friendly and accessible. Examples of these error messages can be found in the packages notebook as well as being provided within the API references (including examples) in the packages [documentation](#).

4.2 Package Structure

The package was structured into two modules, `dual.autodiff` and `autodiff_tools`, to separate the core dual number functionality from the automatic differentiation tools. Despite this the `__init__.py` was configured to expose all core classes and functions to the user from the package's root allowing for intuitive imports, such as:

Listing 7: Example of Package Imports

```
1 from dual_autodiff import Dual, Sin, auto_diff, ...
2 #or
3 from dual_autodiff import *
```

Additionally, the package contains both the Python and Cython versions, `dual.autodiff` and `dual.autodiff_x`, as submodules, allowing users to easily import either version for comparison. Both submodules have been included in the wheel distribution for assessment purposes; however, in practice, only the optimized Cython version would typically be distributed.

Notably, if both versions are imported simultaneously, it can be difficult to distinguish between identical functions, such as `auto_diff`. To avoid this issue, the coursework renames the Cythonised functions on import (typically not required but was included for comparison).

Listing 8: Referencing Cythonised Functions on Import

```
1 from dual_autodiff import auto_diff
2 from dual_autodiff_x import auto_diff as auto_diff_x
```

Note, that when wheel files are used, the user may run into conflicts between the installed and local version of `dual.autodiff_x` in the terminal - discussed in Section 4.4.1.

4.3 Package Distribution and Wheels

The package includes wheels for linux architecture running either python 3.10 and 3.11 in the `wheelhouse` directory (`cp310-manylinux_x86_64` and `cp311-manylinux_x86_64`). These allows the binary distribution of the package, containing both `dual.autodiff` and `dual.autodiff_x` sub-packages. User can thus install the package without performing the cythonisation and compilation locally. Note, the python version would not be typically included, but was included for assessment purposes.

This package also supports local installation (`pip install -e .`), which generates and compiles the current architecture's necessary files locally, such as `.so` files (shared object files). Thus these files are not included outside of the wheels, as they are unique to the users architecture.

4.4 Package dependencies

The package defines both essential or optional dependencies, all of which are defined within `pyproject.toml` file for easy installation. This distinction was made to minimise the dependencies installed on the users device and allow only the required dependencies to be installed. The classes of dependencies are defined in Table 6. Installing these dependencies with the wheels is equally trivial and outlined in the `README` and documentation.

Name	Purpose	Installation
Essential	Core dependencies	<code>pip install -e .</code>
Tutorial	For example notebooks	<code>pip install -e ".[tutorial]"</code>
Testing	For running test suites	<code>pip install -e ".[testing]"</code>
Documentation	Building documentation	<code>pip install -e ".[docs]"</code>

Table 6: Package Dependencies

4.4.1 Neiche Conflict for Coursework

A potential conflict arises when the user installs the package from the wheels but attempts to run the cythonised version from the terminal in the package's directory. In such cases, the locally available package directory is prioritised over the installed version, leading to the program searching for the `.so` files in the local directory, which are only generated during a local installation using `pip install -e ..`. This problem is unique to the coursework as wheels are normally distributed independently to the source code.

4.5 Testing

A comprehensive test suite of all functionalities and edge cases was implemented in the `tests` directory, using `pytest`. This includes 43 tests encompassing 258 statements over 7 files, each focused on testing a specific feature. This allows the user to verify the package's installation, and also stop future developers from introducing conflicts. This was introduced into the continuous integration as a pre-commit hook which is discussed in Section 4.8.

A report on the test coverage was generated using the `coverage` utility, showing a coverage of 100% for the package, shown in Table 7. Note the 82% coverage for the `version.py` file is neglected as it simply contains metadata such as the version number rather than core features.

Name	Stmts	Miss	Cover	Missing
dual_autodiff/__init__.py	5	0	100%	-
dual_autodiff/autodiff_tools.py	103	0	100%	-
dual_autodiff/dual.py	139	0	100%	-
dual_autodiff/version.py	11	2	82%	5-6
TOTAL	258	2	99%	-

Table 7: Test coverage report for the `dual_autodiff` package.

4.6 Documentation

[Documentation](#) on the package was auto-generated using `sphinx`, and is hosted on `readthedocs` which relies on the `.readthedocs.yml` file. This provides a comprehensive guide to the package's functionality, including the API references, examples and context. The documentation can also be built locally using the `make html` command in the `docs` directory, once the `docs` dependencies have been installed.

4.7 Version Control

The package was developed using Git for automatic version control and exploiting `setuptools_scm` for dynamic tracking. Once the first usable version [1.1.0] was released, new features were developed on separate feature branches and merged into the main branch once completed, thus not compromising the package's functionalities.

4.8 Continuous Integration

This package made use of the following continuous integration (CI) methods:

- **Pre-commit Hooks for Testing:** `pytest` was configured to ensure all tests were passed before allowing commits. Ensuring updates would not corrupt the packages. Future developers are encouraged to do the same before submitting a pull request, highlight in the `README`.
- **Synchronisation of Cython:** The `setup.py` was configured to automatically duplicate `.py` files to `.pyx` files before running the build, ensuring the cythonised package was always up to date.
- **Documentation:** `Readthedocs` was configured to automatically build the documentation on each push to the main branch, ensuring the documentation was always up to date.

5 Summary

Overall, the `dual_autodiff` package provides a framework to perform automatic differentiation, supporting a wide range of functions. The package was developed following good software development practices, and shows strong advantages for the accuracy of automatic differentiation over numerical methods. Future versions aim to further exploit C's optimised performance to improve the package's vectorisation.

5.1 Declaration of Use of Autogeneration Tools

This project made use of Large Language Models (LLMs), primarily ChatGPT and Co-Pilot, to assist in the development of the package. These tools have been employed for:

- Helping generate detailed docstrings for the repository’s documentation.
- Formatting plots to enhance presentation quality.
- Performing iterative implementation to predefined code.
- Debugging code and identifying issues in implementation.
- Developing tests for the package.
- Helping with Latex formatting for the report.
- Generating the projects metadata files

References

- [1] Michael Bartholomew-Biggs et al. “Automatic differentiation of algorithms”. In: *Journal of Computational and Applied Mathematics* 124.1 (2000). Numerical Analysis 2000. Vol. IV: Optimization and Nonlinear Equations, pp. 171–190. ISSN: 0377-0427. DOI: [https://doi.org/10.1016/S0377-0427\(00\)00422-2](https://doi.org/10.1016/S0377-0427(00)00422-2). URL: <https://www.sciencedirect.com/science/article/pii/S0377042700004222>.
- [2] Atilim Gunes Baydin et al. *Automatic differentiation in machine learning: a survey*. 2018. arXiv: [1502.05767](https://arxiv.org/abs/1502.05767) [cs.SC]. URL: <https://arxiv.org/abs/1502.05767>.
- [3] H. H. Cheng. “Programming with Dual Numbers and its Applications in Mechanisms Design”. In: *Engineering with Computers, An International Journal for Computer-Aided Mechanical and Structural Engineering* 10.4 (1994), pp. 212–229. URL: https://iel.ucdavis.edu/publication/journal/j_EC1.pdf.
- [4] Annie Cuyt et al. “A Remarkable Example of Catastrophic Cancellation Unraveled”. In: *Computing* 66.3 (May 2001), pp. 309–320. ISSN: 1436-5057. DOI: [10.1007/s006070170028](https://doi.org/10.1007/s006070170028). URL: <https://doi.org/10.1007/s006070170028>.
- [5] William S. Moses et al. “Scalable Automatic Differentiation of Multiple Parallel Paradigms through Compiler Augmentation”. In: *SC22: International Conference for High Performance Computing, Networking, Storage and Analysis*. 2022, pp. 1–18. DOI: [10.1109/SC41404.2022.00065](https://doi.org/10.1109/SC41404.2022.00065).